

Scientific Machine Learning

Recap Lecture

Table of Content

- Linear Models
- Optimization
- Classic ML
- Deep Learning

Linear Models

- Supervised Learning approaches
 - Requires labeled data
- Examples of it are
 - Linear Regression
 - Logistic Regression

$$\left\{ x^{(i)}, y^{(i)} \right\}_{i=1, \dots, m'}$$

Modus Operandi of Linear Regression

1. Formulate a hypothesis for the relation between input x , and output y
2. Look for a strategy to fit the hypothesis parameters, such as gradient descent

$$h(x) = \theta^\top x$$
$$\theta = (X^\top X)^{-1} X^\top Y$$

Probabilistic Interpretation

- Modeling all data points as independent, identically distributed (iid) random variables

$$y^{(i)} = \theta^\top x^{(i)} + \varepsilon^{(i)}$$

- This then allows us to use maximum likelihood estimation.
 1. Construct the log-likelihood $l(\theta)$
 2. Maximize the log-likelihood $\theta = \operatorname{argmax} l(\theta)$

Classification: Logistic Regression

- Apply the sigmoid function on top of the linear regression, and use a different cost function

$$h(x) = \varphi(\theta^\top x) = \frac{1}{1 + e^{-\theta^\top x}}$$

- To apply the maximum likelihood approach, one then needs to construct discrete probability events

$$p(y|x; \theta) = h^y(x)(1 - h(x))^{1-y}$$

Exponential Family

- Consists of Bernoulli, Gaussian, Dirichlet, Gamma, Poisson, and Beta-distribution

$$p(x|\eta) = h(x) \exp(\eta^\top t(x) - a(\eta))$$

- Natural parameter η
- Sufficient statistics $t(x)$
- Probability support $h(x)$
- Log normalizer $a(\eta)$ s.t. probability mass sums to 1

Bayesian Inference: Core Principles

- Uncertainty over parameters
- Probability over the parameter expresses a degree of belief in the result
- Inference over parameters
- Prior knowledge and newly observed data are being combined according to **Bayes Theorem**

$$\mathbb{P}(A|B) = \frac{\mathbb{P}(B|A)\mathbb{P}(A)}{\mathbb{P}(B)}$$

Bayesian Inference

- Seek to arrive at the **posterior distribution**
- Computationally very expensive
 - Unscaled posterior is easy to find, scaled posterior very expensive to find
 - Requires expensive inference approaches such as Markov-Chain Monte Carlo

Bayesian Inference: Monte Carlo Sampling

- Draws from the posterior to find new samples
- Guaranteed to converge if given infinite computation (Monte Carlo theorem)
- In its simplest form

$$\mathbb{E}(h(x)) = \int h(x)p(x)dx$$

Bayesian Inference: Advantages

- Gives us an approximation which can be made as precise as desired through the regulation of the computational budget
- Have probability/certainty estimates over our parameters
 - Always need to consider whether our distributions are converged yet
- Can analyze uncertainties/sensitivities of our model

Sampling Approaches

- Acceptance-Rejection sampling:
$$M \times g_0(\theta) \leq g(\theta)f(y|\theta)$$
- Sampling-importance resampling then extends upon this by introducing a further step in which one resamples from the space of parameters
- Adaptive rejection sampling as the last alternative in case the other two approaches failed

Bayesian Inference: Summary

- Unscaled posterior $g(\theta|y) \propto g(\theta)f(y|\theta)$ contains shape information of posterior
- True posterior obtained by normalizing the unscaled posterior which requires the evaluation of an expensive high-dimensional integral
- Integral usually evaluated numerically using Monte-Carlo approaches

Bayesian Machine Learning

1. Introduces probability distribution over the parameters
2. Perform inference to obtain the probability distribution over the parameters
 - Loss of computational efficiency
 - Gain of uncertainties, especially useful in the small data limit

Optimization

- Optimization seek to solve the minimization problem of the cost function

$$\operatorname{argmin}_{\theta} J(\theta)$$

- If our cost function is convex, then it is easier to optimize
 - Least mean squares loss for linear regression
 - Log-likelihood for logistic regression

Loss functions

- Regression loss
 - L1 loss: $\frac{1}{m} \sum_{i=1}^m |y - h_{\theta}(x)|$
 - MSE loss: $\frac{1}{m} \sum_{i=1}^m (y - h_{\theta}(x))^2$
- Classification loss
 - Cross-entropy loss

$$- \sum_{i=1}^m \sum_{k=1}^K (y_{i_k} \cdot \log h_{\theta,k}(x_i))$$

Gradient Methods

- Gradient descent in its purest form
- Stochastic gradient descent more commonly used
 - Adam
 - Stochastic Gradient Descent

Gradient Methods: Adam

- Adam utilizes the first, and second momentum of the gradient

$$v_t \leftarrow \beta_1 v_{t-1} + (1 - \beta_1) g_t$$

$$s_t \leftarrow \beta_2 s_{t-1} + (1 - \beta_2) g_t^2$$

- typical initialization: $\beta_1 = 0.9, \beta_2 = 0.999$

Gradient Methods: SGD

- Consider the gradient as an expectation, we construct an unbiased estimator of the gradient

$$x \leftarrow x - \eta \nabla f_i(x)$$

- Estimator then constructed as follows

$$\mathbb{E}_i \nabla f_i(x) = \frac{1}{n} \sum_{i=1}^n \nabla f_i(x)$$

- This can then be further enhanced by minibatching

Second-Order Methods

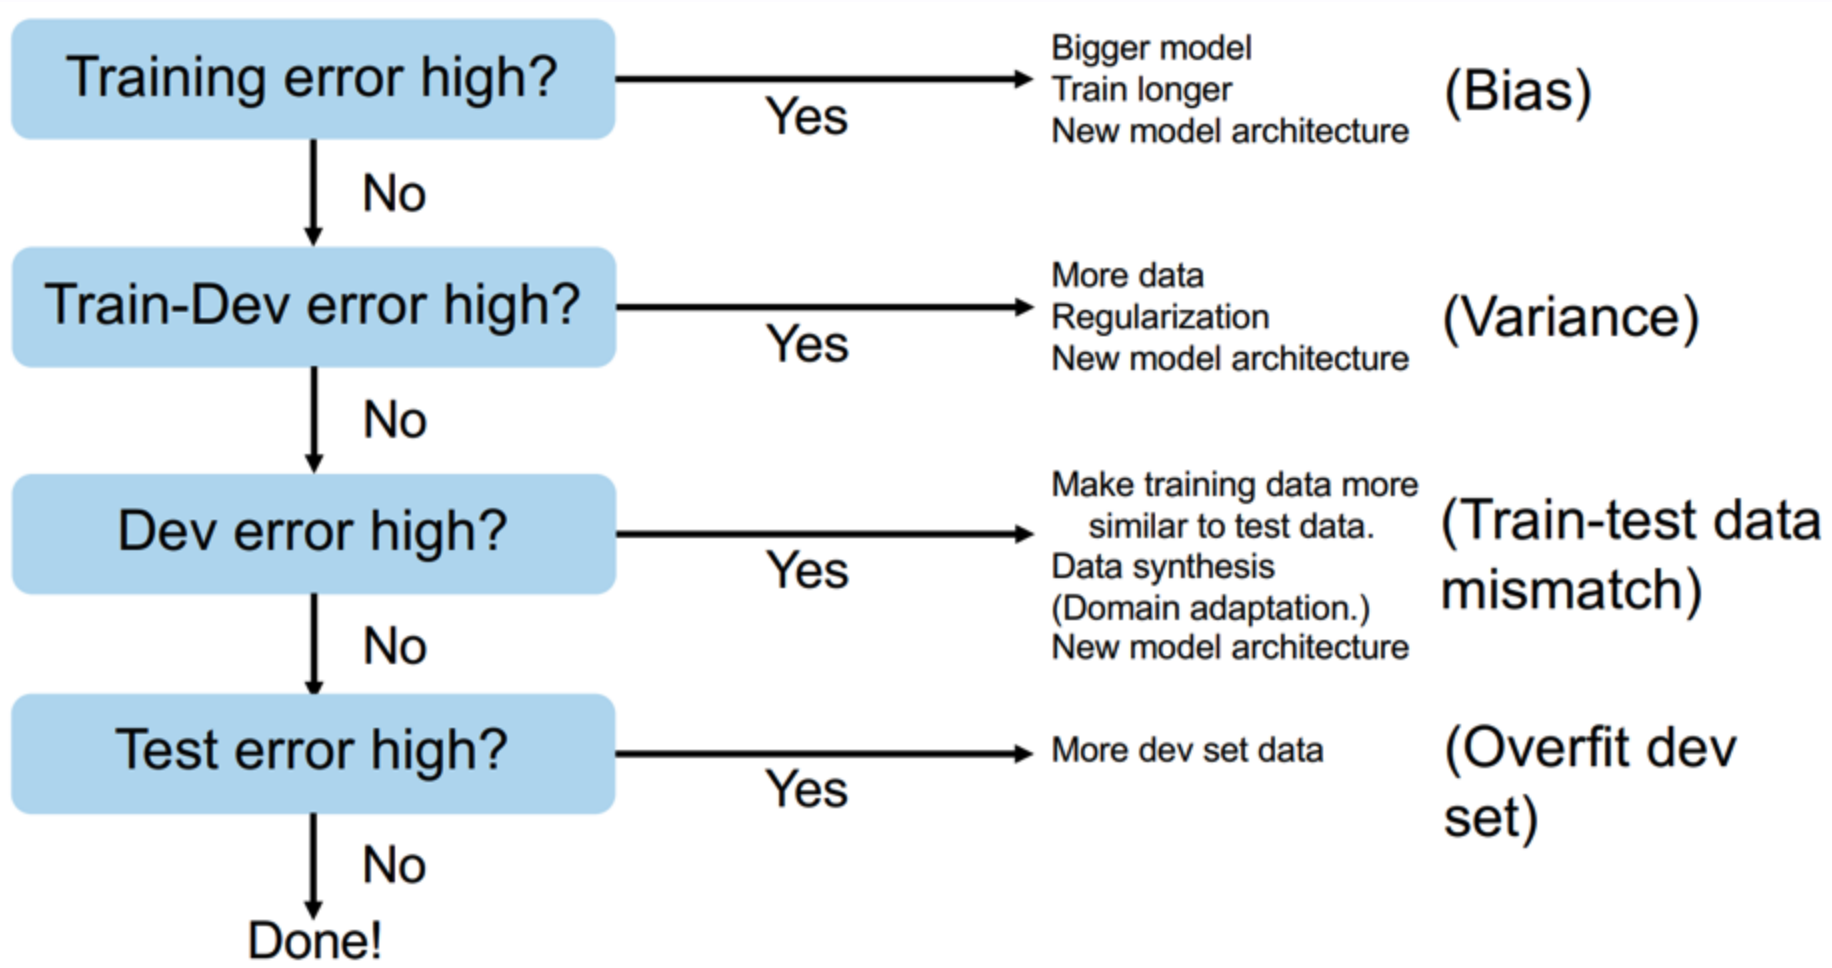
- Don't only take the gradient into account, but also the gradient curvature
- Examples of second order methods include, but are not limited to:
 - Newton's method
 - Quasi-Newton Approach
 - BFGS
 - L-BFGS

Blackbox or Derivative Free Optimization

- Mostly when one cannot compute the gradient in an acceptable timeframe, and can be broken down into the following decision tree
 - Expensive Function
 - Bayesian Optimization
 - Cheap function
 - Stochastic local search
 - Evolutionary search

Tricks of Optimization

- Tricks of Optimization we saw in the lecture are:
 - Data Splitting
 - Regularization
 - Input Normalization, and Parameter Initialization
 - Hyperparameter Search
 - Learning Rate Scheduling



Classical Machine Learning

- Support Vector Machines
- Gaussian Processes

“

jump to lecture script for a quick walk-through

”

Deep Learning

- Gradients (in the presentation)
- Multilayer Perceptron
- Convolutional Neural Networks
- Recurrent Models
- Encoder-Decoder Models

Gradients: The Why

- Writing manual derivatives is tedious, and error-prone
 - Intractable for large code-bases
- Seek to automatically generate the derivatives

Gradients: Forward-Mode Differentiation

- Differentiation performed in the direction of the computational flow with the gradients being accumulated as we go along.
- Does not require elaborate caching, and much less memory

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \left(\frac{\partial c}{\partial b} \left(\frac{\partial b}{\partial a} \frac{\partial a}{\partial x} \right) \right)$$

Gradients: Reverse-Mode Differentiation

- Gradients accumulated in the opposite direction of the computational flow
 - Requires elaborate caching
 - Requires much more memory
 - Results in a much more difficult algorithm

$$\frac{\partial y}{\partial x} = \left(\left(\frac{\partial y}{\partial c} \frac{\partial c}{\partial b} \right) \frac{\partial b}{\partial a} \right) \frac{\partial a}{\partial x}$$

Gradient Forward vs Reverse

- For any function

$$f : \mathbb{R}^n \longrightarrow \mathbb{R}^m$$

- Forward-mode more efficient for gradients of scalar-to-vector functions $m \gg n$
- Reverse-mode more efficient for gradients of vector-to-scalar function $m \ll n$ (most of machine learning)

Multilayer Perceptron

“

[Jump to lecture notes](#)

”